

Utilizing RNN's in 4D Radar Point Clouds for Enhanced Autonomous Vehicle Perception

Nick Timaskovs

K00260158

A Final Year Project submitted in partial fulfilment of the
requirements of the Technological University of the
Shannon for the degree of Bachelor of Science (Honours)
in Software Development.

Supervised By:

William Ward



TUS
Midwest
Department of
Information Technology

May, 2023

Abstract

This thesis explores the multifaceted realm of object detection, tracking, and classification within the dynamic landscape of 4D radar data in autonomous vehicles. The advent of 4D radar technology has opened the way for a new era of high-resolution, real-time data acquisition, enabling an unprecedented understanding of the complex dynamics of objects in motion within a given space and time. This research delves into the development and refinement of algorithms and techniques designed to enhance the accuracy and efficiency of object detection in 4D radar data. By harnessing the power of deep learning models, recurrent neural networks, and advanced signal processing techniques, this study aims to address the challenges posed by noisy and cluttered radar data, while also improving the speed and precision of object detection.

In addition to object detection, this thesis investigates the critical aspect of object tracking in 4D radar data. Tracking objects through space and time is a fundamental task for various applications, including autonomous vehicles, surveillance, and air traffic control. The research proposes novel approaches that leverage both historical and real-time radar data to robustly track objects, accounting for unpredictable movements, occlusions, and changing environmental conditions. Furthermore, this study expands the scope by incorporating classification into the framework, enabling the automatic identification of objects based on their radar signatures. By integrating these elements into a unified system, this thesis contributes to the advancement of 4D radar technology, opening new possibilities for applications in various domains where precise object detection, tracking, and classification is paramount in ensuring safety in autonomous vehicles

Acknowledgements

I want to express my deepest appreciation and gratitude to all those who have supported and guided me throughout the journey of completing my Final Year Project (FYP). This accomplishment would not have been possible without their encouragement, wisdom, and patience.

First and foremost, I would like to extend my sincerest thanks to my FYP supervisor, Mr Willaim Ward, for his unwavering support and guidance throughout this project. His invaluable insights, expertise, and constructive feedback have been instrumental in shaping my work and helping me achieve my goals. I am truly grateful for the time and effort he has invested in mentoring me and the confidence he has shown in my abilities.

Ethical Declaration

I declare that this project and document is wholly my own work except where I have made explicit reference to the work of others. I have read the Department of Information Technology Final Year Project guidelines and relevant institutional regulations, and hereby declare that this document is in line with these requirements.

I have discussed, agreed, and complied with whatever confidentiality or anonymity terms of reference were deemed appropriate by those participating in the research and dealt appropriately with any other ethical matters arising, in line with the TUS Research Ethics Guidelines for Undergraduate and Taught Postgraduate Programmes policy document.

Nick Timaskovs

May 1, 2023

Table of Contents

Abstract	ii
Acknowledgements	3
Ethical Declaration.....	4
Glossary	9
Chapter 1 Introduction	12
1.1 Objective	12
1.2 Project Scope	12
1.3 Solution Developed	12
1.4 Report Structure	13
Chapter 2 Literature Review	14
2.1 Introduction	14
2.2 Autonomous Vehicles	14
2.2.1 Autonomous Vehicle Perception.....	15
2.3 Overview of Radar	16
2.3.1 FMCW Radar Signal Processing.	17
2.4 Overview of LiDAR	19
2.5 Overview of 4D mm Radar	20
2.6 RADAR Vs LiDAR.....	21
2.6.1 4D Radar Advantages.	21
2.6.2 4D RADAR Disadvantages.....	21
2.7 Deep Learning	23
2.7.1 Convolutional Neural Networks (CNNs).....	24
2.7.2 Recurrent Neural Networks (RNNs).....	25
2.7.3 Region-Based Convolutional Neural Network (R-CNN).	26
2.8 Object Detection & Tracking	27
2.8.1 Traditional Object Detection Techniques.	28

2.9	Object Classification	29
2.9.1	AlexNet (2012).....	30
2.9.2	VGGNet (2014).....	30
2.9.3	ResNet (2015).	30
2.9.4	DenseNet (2017).	30
2.10	Related Experimental Results	30
2.11	Conclusion	31
Chapter 3	Analysis and Design.....	32
3.1	Application Overview	32
3.2	Potential Blockers and Considerations	32
3.2.1	Data Sparsity and Quality	32
3.2.2	Dataset Acquisition	32
3.2.3	Computational Complexity	32
3.2.4	Steep Learning Curve.....	33
3.3	Scope	33
3.4	Prerequisites	34
3.5	Software Development Process	34
3.5.1	Agile.....	34
3.5.2	Minimal viable product.....	35
3.5.3	Project Management.....	36
3.6	Tools and Framework Considered	36
3.6.1	Python	36
3.6.2	Jupyter Notebook	37
3.7	Dataset	39
3.7.1	View-Of-Delft Sensor Setup	40
3.7.2	Dataset Frame Information	40
3.7.3	Dataset Label Information.....	40

3.8	Data Visualization and Exploration	41
3.9	Model Architecture Considered	42
3.9.1	PV-RCNN Components Overview	42
3.10	Training and Inference Details	44
3.11	Evaluation Metrics	44
Chapter 4	Implementation	46
4.1	Introduction	46
4.2	Tools Used.....	46
4.2.1	Integrated Development Environment (IDEs)	46
4.2.2	Ubuntu.....	46
4.2.3	Libraries	47
4.3	Source Control.....	48
4.4	Installation	48
4.5	Data Loader & Split	48
4.6	Point Cloud Visualization in Open3d.....	51
4.7	Getting required repositories	53
4.8	GitHub Contributions	53
4.9	Building the model	53
4.10	Training the model.....	53
Chapter 5	Testing and Results	54
5.1	Introduction	54
5.2	Unit Testing.....	54
5.2.1	What is Unit Testing, and why is important?.....	54
5.2.2	Testing Framework in PyCharm	54
5.3	Unit Testing Done	54
5.4	Result Analysis.....	54
5.5	Discussion of Findings	55

Chapter 6	Conclusion:	56
6.1	Future Work	56
Chapter 7	Appendices	57

Glossary

ADAS: Advanced Driver Assistance Systems

AI: Artificial Intelligence

AoA: Angle Of-Arrival

AV: Autonomous Vehicle

CNN: Convolutional Neural Networks

CFAR: Constant False Alarm Rate

CW: Continuous Wave

DBSCAN: Density-Based Spatial Clustering of Applications with Noise

DDM: Doppler-Division Multiplexing

FDM: Frequency-Division Multiplexing

FFT: Fast Fourier Transform

FoV: Field of View

FMCW: Frequency-Modulated Continuous Wave

FPS: Frames Per Second

GPS: Global Positioning System

GPU: Graphics Processing Unit

HOG: Histogram of Oriented Gradients

IDE: Integrated Development Environment

LiDAR: Light Detection and Ranging

LSTM: Long Short-Term Memory

MEMS: Micro-Electromechanical System

MIMO: Multiple Input Multiple Output

ML: Machine Learning

mAP: Mean Average Precision

MVP: Minimal Viable Product

NN: Neural Network

OpenCV: Open-Source Computer Vision Library

PV-RCNN: Point-Voxel Feature Set based 3D Object Detection from Point Clouds

PyTorch: An open-source machine learning library

R-CNN: Region-Based Convolutional Neural Network

RAD: Range-Doppler

RADAR: Radio Detection and Ranging

RAED: Range-Azimuth-Elevation-Doppler

RCS: Radar Cross Section

RD: Range-Doppler

RoI: Region of Interest

RNN: Recurrent Neural Networks

Rx: Receiver antenna

SAE: Society of Automotive Engineers

SLAM: Simultaneous Localization and Mapping

SNR: Signal-to-Noise Ratio

SORT: Simple and Online Real-time Tracking

TDM: Time-Division Multiplexing

Tx: Transmitter antenna

VOD: View-Of-Delft (Dataset)

Table of Figures

Chapter 1 Introduction

This thesis investigates object detection, tracking, and classification using 4D radar in autonomous vehicles, focusing on improving accuracy and efficiency with advanced algorithms, deep learning, and signal processing. It addresses challenges in handling noisy radar data to enhance detection speed and precision.

1.1 Objective

This project aims to develop a web-based time management and user identification system that leverages facial recognition technology. The technology employs cameras to take a picture of the user's face and then compare it to a previously registered image in the system to verify the user's identity. The system grants access when the user's identification has been verified and logs the user's entry and exit times. This method may be utilised in various settings, such as businesses, institutions of higher learning, and healthcare facilities, to enhance security and time management.

This solution's objectives are:

- To overcome the difficulty working with 4D radar data.
- To explore the theoretical utilisation of RNNs for object detection in 4D radar.

The academic objectives:

- To learn about deep learning.
- To learn and understand 4D radar point cloud data structure.
- To learn and understand how object detection, tracking and classification is done in 4D radar.
- Gain experience with academic research.

1.2 Project Scope

Pending.

1.3 Solution Developed

The developed source code is available on github: <https://github.com/nicktmv/four-d-radar-thesis> with all the neccary steps contained in the README.md file.

1.4 Report Structure

This report comprises five main chapters: Literature Review, Analysis and Design, Implementation, Testing and Results and Conclusions.

- Literature Review: This chapter surveys existing research and developments in the field to establish the study's context and foundation.
- Analysis and Design: It outlines the theoretical framework and design methodology used for the study's objectives.
- Implementation: This section describes the practical application of the designed framework and the development of the study's solution.
- Testing and Results: It presents the evaluation methods, testing procedures, and the outcomes of the implemented solution.
- Conclusions: This final chapter summarizes the study's findings, discusses their implications, and suggests areas for future research.

Chapter 2 Literature Review

2.1 Introduction

This chapter explores how various large software companies have tackled the challenge of facial recognition technology and the underlying mechanisms that power it. The author examines these companies' various techniques and methodologies and explains their solutions. The primary objective of this chapter is to equip the reader with comprehensive knowledge of contemporary facial recognition software and the cutting-edge technologies that drive its operation.

2.2 Autonomous Vehicles

Autonomous vehicles (AVs) are self-driving vehicles that are equipped with a multitude of sensors illustrated by **Figure 1**, including LiDAR, radar, cameras, ultrasonic sensors, and GPS. These sensors provide information about the vehicle's surroundings, capturing data about other vehicles, pedestrians, road infrastructure, and weather conditions. AVs and ADAS also use advanced AI and ML algorithms to process sensor data to make real-time decisions and interact with vehicle Control Systems such as steering, acceleration, and braking. (Synopsys.com, 2023)




Figure 1: Sensor setup of the Valeo Drive4U prototype. (Valeo 2018)

2.2.1 Autonomous Vehicle Perception.

Perception is a fundamental component of AVs and advanced driver assistance systems (ADAS), referring to how a vehicle senses or “sees”, enabling them to understand their environment, including identifying and tracking objects, pedestrians, other vehicles, and road infrastructure such as lane markings, traffic signs, and traffic lights. Perception is done with the use of a multitude of sensors such as RADAR, LiDAR, Cameras and Ultrasonic.

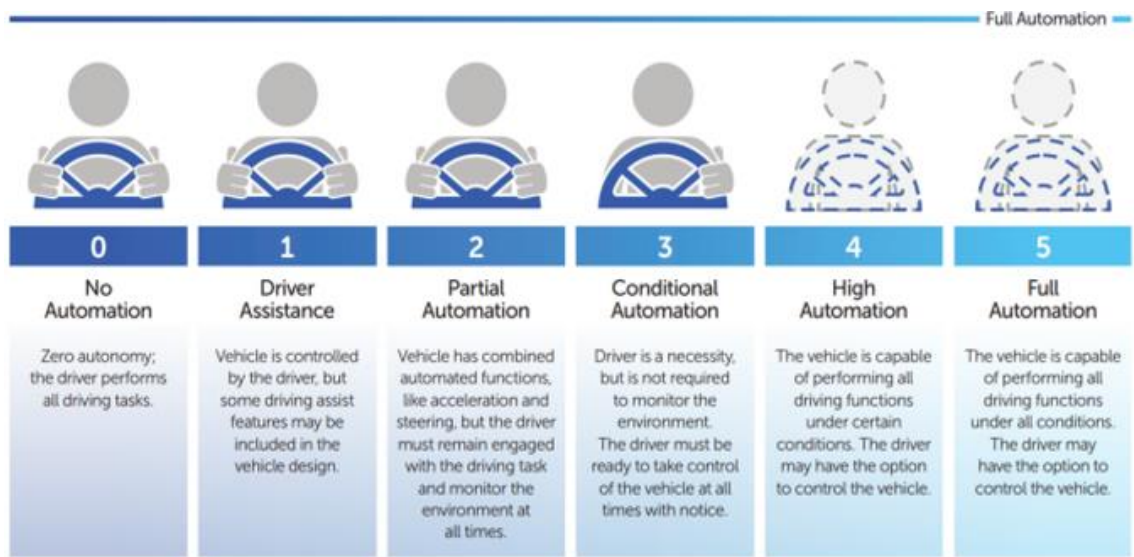


Figure 2: (2.0 A Vision for Safety AUTOMATED DRIVING SYSTEMS, n.d.)

Illustrated by **Figure 2**, AD and ADAS are categorized into 6 levels, Level 0 (full human control) to Level 5 (complete vehicle autonomy) these levels are defined by the Society of Automotive Engineers (SAE) in their J3016 standard, which provide a framework to understand the progression of autonomy in vehicles.

As ADAS research, testing, and implementation in vehicles continue to grow worldwide, there is an increasing focus on establishing standardized rules and regulations to guarantee their secure incorporation into society. Most commercial vehicles are categorized as Level 1 to Level 2 autonomy due to limitations in sensors, cost factors, and the need for ongoing driver attention and control. These vehicles generally come equipped with ADAS features such as emergency braking, blind spot detection, Adaptive Cruise Control, Automatic Parking, and Lane Assist. (2.0 A Vision for Safety AUTOMATED DRIVING SYSTEMS, n.d.)

At CES 2023, Mercedes-Benz made a significant announcement, stating that it had achieved Level 3 (L3) autonomous driving certification in the United States, specifically from the state of Nevada. It's worth noting that L3 certification is granted at the state level in the US, so Mercedes' system is only considered L3 in Nevada for now. This move by Mercedes is expected to encourage other major automotive manufacturers such as Hyundai-Kia, Stellantis, BMW, GM, and Honda to pursue Level 3 autonomous driving technology, as they have also been reporting progress and plans for L3 rollout. (AUTOCRYPT, 2023). SLAM algorithms are employed in autonomous vehicles, they enable the vehicle to create a map of its surroundings while determining its own position within that map simultaneously. SLAM algorithms empower the vehicle to chart unexplored environments, and engineers utilize this map data for tasks like planning routes and avoiding obstacles. (Mathworks.com, 2023)

AVs use sensor fusion techniques to combine data from multiple sensors, improving the accuracy and robustness of perception systems. Kalman filters and Bayesian approaches are commonly employed for sensor fusion. (Anwesh Marwade, 2020)

2.3 Overview of Radar

Radar technology has become the basis for modern sensing and surveillance systems for several decades. Its applications span from military, aviation and maritime navigation to weather monitoring and autonomous vehicles. Radar technology has its roots in the early 20th century, with significant advancements occurring during World War II. Early radar systems used microwave signals to detect and locate objects. Notable developments include the British Chain Home system and the American SCR-270 radar. These systems provided critical advantages in terms of early warning and target tracking. (Bloom, 2020)




Figure 3: Basic Radar Operation (Roshni Y, 2019)

Illustrated by **Figure 3**, Radar operates on the principles of emitting electromagnetic waves, using a transmitter antenna (Tx) and receiving their echoes via a receiver antenna (Rx) after reflecting off objects and analyzing the time delay and **Doppler shift** of these echoes to determine an object's range, speed, and direction. Key components include transmitters, antennas, receivers, and signal processing units.

Doppler Shift Definition: *“the apparent difference between the frequency at which sound or light waves leave a source and that at which they reach an observer, caused by the relative motion of the observer and the wave source”*. ([Doppler effect](#) | Definition, Example, & Facts| Britannica, 2023)

2.3.1 FMCW Radar Signal Processing.

There are many different types of RADAR, all used in various use cases as discussed previously however mentioned in **section [2.3]**, FMCW is the most used in automotive. and is relevant to this use case.

FMCW operates using linear frequency-modulated continuous-wave signals to measure range, angle, and velocity. Based on regulations they can use two frequency bands: 24 GHz and 77 GHz, with a preference for 77 GHz due to its wider bandwidth, higher Doppler resolution, and smaller antennas great for long-range use cases. (Udemy, 2023)

FMCW signals have key parameters as shown in **Figure 4**: start frequency (f_c), sweep bandwidth (B), chirp duration (T_c), and slope (S). A frame consists of chirps with a frame time of (T_f). A frequency mixer combines received and transmitted signals to produce two signals: sum frequency $f_T(t) + f_R(t)$ and difference frequency $f_T(t) - f_R(t)$. with a low-pass filter used to obtain the intermediate frequency (IF) signal. Complex exponential IF signals are achieved in practice using a quadrature mixer. (Zhou et al., 2022)

Range and Doppler velocity can be estimated from the IF signal, leading to a 2D complex data matrix called the Range-Doppler (RD) map. The angle information is obtained using multiple receivers or transmit channels. (Zhou et al., 2022)

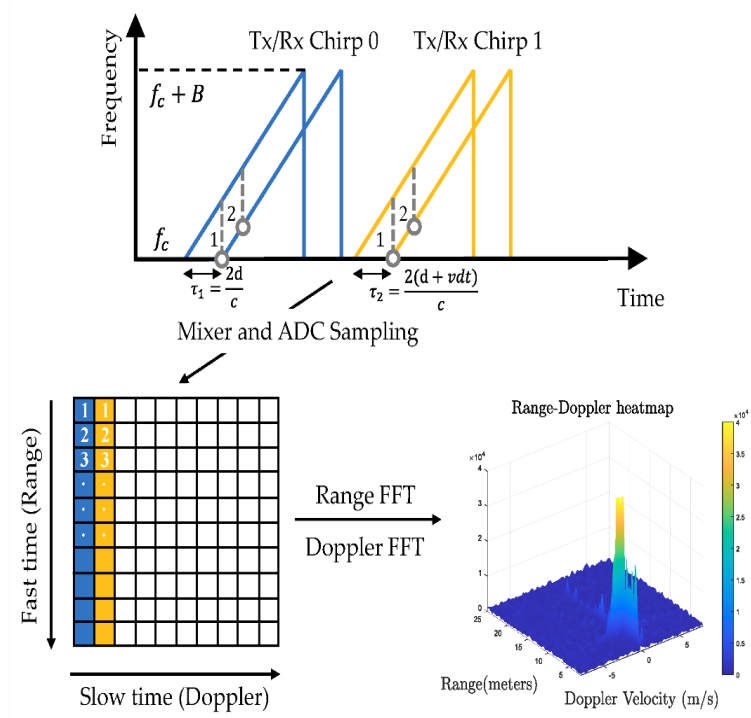


Figure 4: Radar Tx/Rx signals and the range-Doppler map (Zhou and Yue, 2022)

Angle estimation in radar can be achieved using SIMO radar, which involves a single transmitter antenna and multiple receive antennas. By measuring the phase change between adjacent receive antennas, the direction of an object can be calculated using the formula $\Delta\phi = 2\pi d \sin\theta / \lambda$, where θ represents the object's angle, and d is the antenna spacing. (Zhou et al., 2022)

To achieve maximum angle precision, the antenna spacing can be set to $\lambda/2$, and a third Fast Fourier Transform (FFT) is used for processing.

The angular resolution of a SIMO radar depends on the number of receive antennas, but this number is limited by the cost of signal processing.

MIMO radar employs multiple transmit and receive channels, creating a virtual array with many channels. Various techniques like time-division multiplexing (TDM), frequency-division multiplexing (FDM), and Doppler-division multiplexing (DDM) are used to ensure signal separation.

TDM is straightforward, where different transmit antennas take turns sending signals. DDM sends all signals simultaneously but reduces Doppler velocity resolution.

Once signals are separated, a 3D tensor, called the RAD tensor, is created by stacking Range-Doppler (RD) maps. Angular resolution can be enhanced using super-resolution techniques like Capon, MUSIC, and ESPRIT.

The radar detection process involves coherent integration to improve signal-to-noise ratio (SNR), followed by a constant false alarm rate (CFAR) detector for peak detection. Angle estimation is then applied, resulting in a point cloud with range, Doppler, and angle measurements.

In conventional radars, only azimuth angles are resolved, while 4D radars output both azimuth and elevation angles.

Low CFAR thresholds are used in safety-critical applications for high recall. Spatial-temporal filtering techniques, such as DBSCAN and Kalman filtering, are employed to reduce errors caused by clutter and interference.

2.4 Overview of LiDAR

LiDAR, an abbreviation for "Light Detection and Ranging," constitutes a remote sensing method utilizing laser light for precise distance measurements. This technique facilitates the generation of intricate 3D maps of surroundings. LiDAR functions by emitting quick laser pulses and precisely gauging the duration for the reflected light to travel back.

LiDAR functions through a typical system comprising three fundamental elements:

Laser Source: Utilizing lasers, LiDAR emits rapid bursts of light, usually within the near-infrared spectrum. These laser pulses travel through the atmosphere until encountering objects such as terrain, trees, or structures.

Scanner: To ensure comprehensive data collection, a scanner, often a rotating mirror or micro-electromechanical system (MEMS), directs the laser pulses in various directions.

Detector: An incredibly sensitive instrument, the detector measures the time taken for the laser pulses to reach objects and return to the LiDAR system. This principle of "time-of-flight," in conjunction with the constant speed of light, facilitates precise distance calculations.

When a laser pulse strikes an object like a tree or a structure, the light reflects to the LiDAR system. The scanner, typically a rotating mirror or MEMS, captures and directs the returning light to the detector. The detector, being highly sensitive, measures the time it takes for light to travel back and forth between the LiDAR system and the object. This principle of time-of-flight, combined with the constant speed of light, enables LiDAR to accurately compute the distance to an object.

2.5 Overview of 4D mm Radar

4D-imaging radar represents a cutting-edge iteration of mm Wave radar technology that surpasses conventional radar's capabilities. It deploys echolocation and utilizes time-of-flight measurement principles to create a representation of objects within a 3D setting. The four dimensions include **Range, Azimuth, Elevation, and Doppler velocity**. It also provides some other low-level capabilities such as radar-cross-section (RCS) or signal-to-noise ratio (SNR). (Everythingrf.com, 2021)

To create a detailed representation of the surrounding environment for a vehicle, a 4D imaging radar employs a Multiple Input Multiple Output (MIMO) antenna array. This array can consist of numerous antennas that transmit signals towards objects in the vicinity of the device and subsequently collect the reflected signals. The information gathered by these antennas is then utilized to construct a point cloud, depicting the region encompassing the vehicle in high detail. This point cloud captures not only the location and movement of objects but also their shape and texture, allowing for an unprecedented level of environmental awareness. By analysing the Doppler shifts in the returned signals, the system can determine the velocity of moving objects, adding a dynamic aspect to the static scene captured by traditional radar systems. The use of a MIMO antenna array significantly enhances the radar's resolution and accuracy, enabling the detection of smaller objects and finer details in complex urban and highway environments. Consequently, 4D imaging radar plays a crucial role in enhancing the safety and efficiency of autonomous driving systems, offering a comprehensive understanding of the vehicle's immediate surroundings and potential hazards.

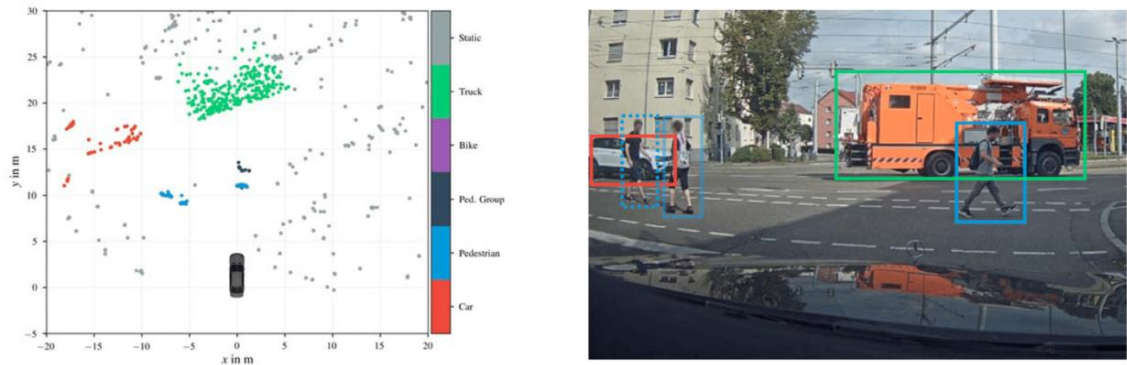


Figure 5: Examples of 4D RADAR objects detected projected to the image plane.
(Zhou et al., 2020)

2.6 RADAR Vs LiDAR

2.6.1 4D Radar Advantages.

Unlike RGB cameras, which utilize the visible light spectrum (384-769 terahertz), and LiDAR systems, which operate in the infrared spectrum (361-331 terahertz), Radars operate at significantly longer radio wavelengths (77-81 gigahertz). This characteristic allows Radars to provide reliable measurements (Cohen, 2023) even in adverse weather conditions such as rain, fog, or snow, ensuring reliable performance in various environments. (Yang et al., n.d.)

RADAR operates independently of ambient lighting conditions, making it ideal for 24/7 surveillance and autonomous driving at night (Zhou et al., 2020). By breaking down each azimuth into an array of intensity values distributed radially, RADAR introduces an additional dimension that LiDAR lacks. This unique feature enables RADAR to construct a top-down, image resembling a photograph, a task that a LiDAR unit cannot accomplish without incorporating multiple channels. (Navtech Radar, 2023)

4D RADAR offers notable advantages in terms of long-range detection capabilities, with a range extending beyond 200 meters, as well as robustness. Millimetre-wave RADAR can penetrate certain non-metallic obstacles, such as plastic and fabric allow for RADAR to be seamlessly hidden behind a bumper for an aesthetic look. (Xx, Xx and Xxxx, n.d.)

2.6.2 4D RADAR Disadvantages.

By integrating data from various sensors, such as visual sensors and LIDAR, the system can enhance its understanding of the driving environment. Deep learning (DL) techniques

have played a pivotal role in this, with a multitude of advanced deep neural networks (DNNs) being employed for perception tasks. Thanks to the substantial learning capacity of DL, there has been a significant improvement in the performance of these tasks. Many DL frameworks have been explored for processing both image and LIDAR data, as they provide ample data for training and validating deep neural networks. In contrast, research on RADAR-related DL studies has been limited, primarily due to the sparsity of RADAR data. (Zhou et al., 2020). For more research to be completed using RADAR more data is required. LiDAR excels in providing higher accuracy by being able to output over 100,000 points per frame while 3D RADAR outputs only 1,000 points per frame. (Hesai Webmaster, 2023)

Illustrated by **figure 6**, certain objects that are clearly visible in LiDAR data, such as cars and pedestrians, may appear blurred or less well-defined when observed through RADAR data. (AR, 2021)




Figure 6: LiDAR VS RADAR Point Cloud (AR, 2021)

To conclude **Figure 7:** summarises some of the advantages of RADR over LiDAR, however because of Radars limited resolution and the absence of semantic features, RADAR-based technologies “*for detection and tracking, vehicle self-localization, and HD map updating currently*” not as advanced as other perception sensors in highly autonomous driving. Nevertheless, research efforts in RADAR technology have been on the rise, thanks to the unique advantages offered by RADAR sensors. Enhancing the quality and imaging capabilities of millimetre-wave (MMW) RADAR data, along with exploring the full potential of RADAR sensors, is essential for gaining a comprehensive understanding of the driving environment. (Zhou et al., 2020)

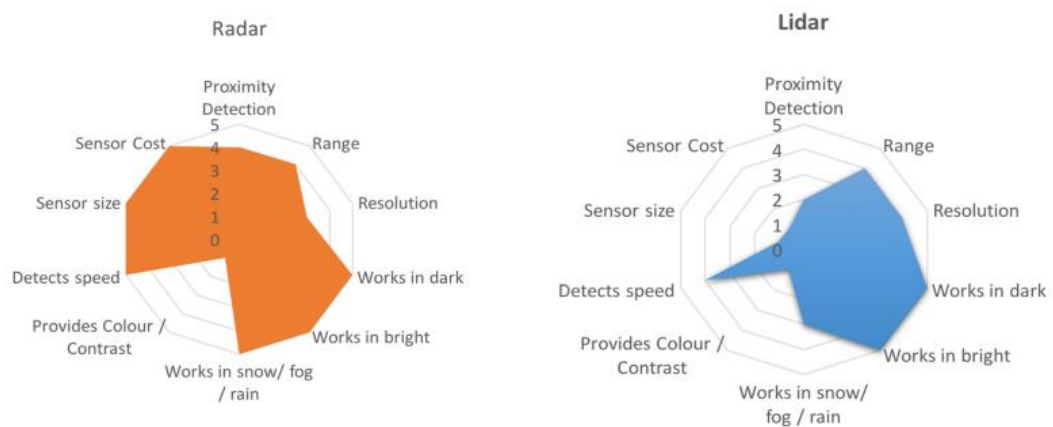


Figure 7: RADAR VS LiDAR Graph (Barnard, 2016)

2.7 Deep Learning

Deep learning models consist of a series of interconnected layers, like the human brain that consist of millions of interconnected neurons that cooperate to acquire knowledge. (Gillis, Burns and Brush, 2023) these layers create a continuous and trainable model using backpropagation. Researchers have extensively studied these layers and neural network structures in recent years, leading to enhancements in data feature representations. The utilization of extensive annotated datasets has enabled the training of models with larger parameter counts, resulting in improved performance on public benchmarks and challenges.

Deep Learning applications employ a hierarchical set of algorithms referred to as an artificial neural network (ANN). The architecture of such an ANN is inspired by the

biological neural networks found in the human brain, enabling a learning process that surpasses the capabilities of conventional machine learning models. (Levity.ai, 2023)

This section will focus on the state-of-the-art deep learning models for computer vision tasks such as object detection, tracking and classification.

2.7.1 Convolutional Neural Networks (CNNs).

Images captured by cameras are represented as 3D matrices with pixel values ranging from 0 to 255 for each of the Red, Green, and Blue (RGB) channels. When utilizing images as input for a machine learning model, it becomes computationally intensive. For instance, if we consider an image with dimensions 200x200x3, a single fully connected layer requires 120,000 parameters to process each pixel. This conventional ANN cannot efficiently handle high-dimensional input data and cannot directly learn spatial features from the input. (Ouaknine, 2022)

In recent years, Convolutional Neural Networks (CNNs) have been explored to address these challenges. CNNs utilize Convolutional layers are comprised of a grid of neurons, with a requirement that the previous layer also consists of a grid of neurons in a rectangular shape. Each neuron in this layer receives inputs from a rectangular segment of the preceding layer, with the same set of weights applied to this rectangular segment for all neurons within the convolutional layer. Consequently, the convolutional layer essentially performs an image convolution operation on the previous layer, using these weight values to define the convolution filter. (Joel Markus Vaz and Balaji, 2021)

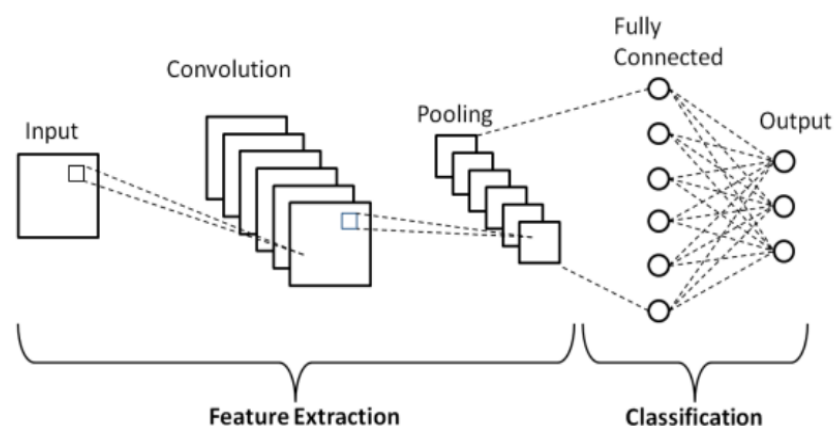


Figure: 8 Diagram of a CNN (ResearchGate, 2019)

Within each convolutional layer, there can be multiple grids, each of which obtain inputs from all the grids in the preceding layer, utilizing potentially distinct filters. Following each convolutional layer, there is a pooling layer illustrated by **Figure 8**.

The pooling layer selects small rectangular blocks from the convolutional layer and reduces them to a single output from that block through subsampling. Various pooling methods can be employed, such as averaging, taking the maximum value, or using a learned linear combination of the neurons within the block. In this context, our pooling layers consistently utilize max pooling, which means they pick the maximum value from the block they are pooling. (Joel Markus Vaz and Balaji, 2021)

Finally, after a series of convolutional and max-pooling layers, the neural network's high-level reasoning occurs through fully connected layers. A fully connected layer connects all the neurons from the previous layer, whether it was a fully connected, pooling, or convolutional layer, to every neuron within itself. Fully connected layers do not have a spatial arrangement (they can be envisioned as one-dimensional), making it infeasible to introduce convolutional layers after a fully connected layer. (Gibiansky, 2014)

2.7.2 Recurrent Neural Networks (RNNs).

Recurrent neural networks (RNNs), like other deep learning methods, have been around since the 1980s, but their true capabilities became evident only recently. The introduction of long short-term memory (LSTM) in the 1990s, along with greater computing power and the abundance of data, has propelled RNNs to the forefront of machine learning. (Kalita, 2022)

An RNN is a neural network designed for handling sequential data, and it finds application in various fields, including temporal series analysis, such as music, video, and stock market data, as well as Natural Language Processing tasks like textual analysis and translation.

In an RNN, the output from the previous step serves as input to the current step. Unlike traditional neural networks, where inputs and outputs are treated independently. (GeeksforGeeks, 2018)




Figure: 9 RNN Architecture (Avijeet Biswal, 2020)

Illustrated by **Figure: 9** "x" represents the input layer, "h" corresponds to the hidden layer, and "y" denotes the output layer. The network's performance is enhanced by utilizing parameters A, B, and C. At each time step "t," the present input is a composite of both the current input at $x(t)$ and the previous input at $x(t-1)$. The output at any given time is looped back into the network to refine the model's output. (Avijeet Biswal, 2020)

Recurrent Neural Network (RNN) models offer several advantages, such as their ability to handle sequences of varying lengths, their consistent number of parameters regardless of input size, and their capacity to incorporate historical context by sharing parameters across timestamps. Nevertheless, RNNs are hindered by their slow training speed and the issue of vanishing gradients, which tend to occur towards the end of sequences. Additionally, these models struggle to effectively capture long-term dependencies due to information loss during sequence processing. (Ouaknine, 2022)

2.7.3 Region-Based Convolutional Neural Network (R-CNN).

This approach begins by employing a selective search technique. Initially, it segments an image into smaller regions and then merges them hierarchically using various colour spaces and similarity metrics (Jasper et al., 2013). This process yields a limited set of region proposals that may potentially contain objects of interest. The R-CNN model, proposed by (Girshick et al., 2016), combines this selective search method for region proposal generation with deep learning for object classification illustrated by **Figure 10**.




Figure: 10 Diagram of a RCNN (Gandhi, 2018)

Each region proposal is resized to match the input requirements of a Convolutional Neural Network (CNN), which produces a 4096-dimensional feature vector. This feature vector is then used as input for binary Support Vector Machine (SVM) classifiers, one for each class. Additionally, it is used in a linear regressor to adapt the shapes of the corresponding bounding boxes, reducing location errors. (Hearst et al., 1998)

The CNN is trained on the 2012 ImageNet dataset for image classification and then fine-tuned using region proposals that have an Intersection over Union (IoU) greater than 0.5 with the ground-truth bounding boxes. Two versions of the model are created: one using the 2012 PASCAL VOC dataset and the other using the 2013 ImageNet dataset with associated bounding boxes. SVM classifiers are also trained for each class within both datasets. The top-performing R-CNN models have achieved a mean Average Precision (mAP) score of 62.4% in the 2012 PASCAL VOC challenge, representing a substantial 22.0-point improvement over the second-best result on the leaderboard. Furthermore, they achieved a 31.4% mAP score on the 2013 ImageNet dataset, surpassing the second-best result on the leaderboard by 7.1 points. (Ouaknine, 2022)

2.8 Object Detection & Tracking

Object detection represents a fundamental task in the realm of computer vision and holds significant importance in enabling autonomous driving. Autonomous vehicles heavily rely on their ability to perceive their surroundings, ensuring safe driving performance. To achieve this, they utilize object detection algorithms, which accurately identify objects such as pedestrians, vehicles, traffic signs, and barriers within their proximity. Deep

learning-based object detectors are instrumental in real-time identification and localization of these objects. (Balasubramaniam and Pasricha, n.d.)

This section explores the current state of the art in object detection and highlights the open challenges associated with integrating these technologies into autonomous vehicles.

2.8.1 Traditional Object Detection Techniques.

The modern progression of object detection techniques initiated two decades ago with the Viola Jones detector, in 2001, Paul Viola and Michael Jones introduced an object recognition framework for real-time human face detection. This framework employs sliding windows to scan an image at various locations and scales to identify human faces. The search is based on "haar-like" features, named after Alfred Haar, who pioneered the concept of haar wavelets. These wavelets serve as the image's feature representation. To expedite detection, an integral image is utilized, ensuring that the computational effort for each sliding window remains independent of its size. (baeldung, 2022).

The authors also employ the Adaboost algorithm for feature selection, which identifies a small set of features that are particularly useful for face detection from a large pool of random features. Additionally, the algorithm incorporates Detection Cascades, a multi-stage detection approach aimed at reducing computational overhead. This means that it prioritizes less computation on background windows and focuses more on potential face targets. (Borah, 2020)

A few years later, the Histogram of Oriented Gradient (HOG) detectors gained popularity, especially for detecting pedestrians (Tyagi, 2021). These HOG detectors were subsequently expanded into Deformable Part-based Models (DPMs), representing the first models geared towards the detection of multiple objects. (Davies, 2022)

Around 2014, the surge in interest surrounding deep neural networks led to a significant breakthrough in multiple object detection with the introduction of the Regions with Convolutional Neural Network (R-CNN) deep neural network model. This innovation resulted in a remarkable 95.84% enhancement in Mean Average Precision (mAP) over the existing state-of-the-art methods.

This pivotal advancement not only redefined the effectiveness of object detectors but also made them appealing for entirely new application domains, particularly in the context of Autonomous Vehicles (AVs).

Since 2014, the continued evolution of deep neural networks and the progress in Graphics Processing Unit (GPU) technology have paved the way for faster and more efficient object detection in real-time images and videos. Modern AVs heavily rely on these improved object detectors for various crucial tasks, including perception, path planning, and other decision-making processes. **Figure 11** illustrates the taxonomy of object detectors.

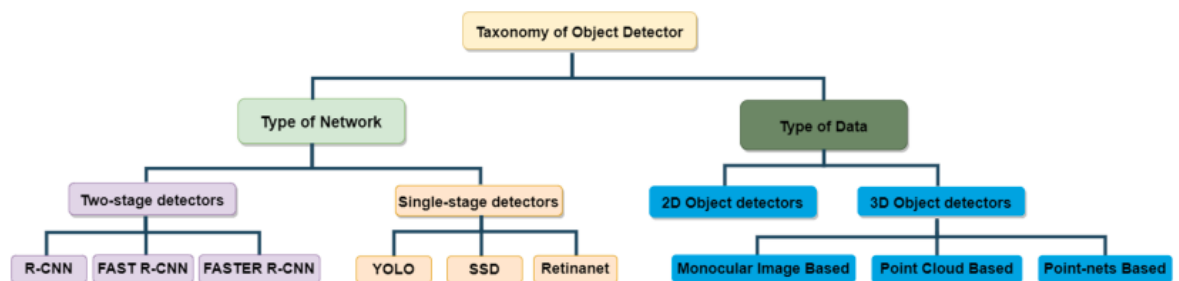


Figure 11: Taxonomy of object detectors (Balasubramaniam and Pasricha, 2022)

2.9 Object Classification

One of the widely pursued objectives in the field of computer vision involves the task of classification, which entails assigning a specific category to every image within a dataset. The ImageNet dataset, introduced by (Deng et al., 2009), has played a significant role in extensive exploration and research on classification. A vector is used to measure and represent the local characteristics of an image, summarizing these features in the form of a histogram that describes the overall image. This section will provide an overview of the evolution and improvements made over the years in image classification.

This section discusses the accomplishments in deep learning related to image classification, specifically in the context of the ImageNet challenge. It provides detailed information about well-known modules and architectures that are still widely employed for feature extraction today. It is important to note that this isn't an exhaustive catalogue of all the models developed between 2012 and 2018. Furthermore, the more recent

advancements, particularly in Transformer architectures as proposed by (Dosovitskiy et al., 2020) fall outside the scope of this thesis.

2.9.1 AlexNet (2012).

A deep convolutional neural network known as AlexNet, created by Alex Krizhevsky and his team in 2012, emerged victorious in the ImageNet Large Scale Visual Recognition Challenge (ILSVRC). This groundbreaking architecture featured numerous convolutional layers and served as a pioneering example of how deep learning could be harnessed for the purpose of image classification. (Nitin Kushwaha, 2023)

2.9.2 VGGNet (2014).

The VGGNet design, created by the Visual Geometry Group at the University of Oxford, focused on creating deeper neural network architectures. Its straightforward and consistent structure set a standard for investigating network depth. (Deepchecks, 2021)

2.9.3 ResNet (2015).

Residual Networks, commonly known as ResNets, revolutionized deep neural networks by introducing the concept of residual connections, effectively tackling the issue of vanishing gradients. This breakthrough enabled the successful training of extremely deep networks, leading to the practical realization of ResNets with hundreds of layers, resulting in significant performance enhancements. They achieved remarkable success by securing the top position in the ILSVRC 2015 classification competition with a top-5 error rate of 3.57%, using an ensemble model. (Great Learning Team, 2020)

2.9.4 DenseNet (2017).

DenseNet introduced a novel approach to connectivity patterns by promoting dense connections, which encouraged the reutilization of features and facilitated the flow of gradients. This architectural innovation showcased enhanced training effectiveness and increased accuracy in the context of image classification tasks. (Paperswithcode.com, 2020)

2.10 Related Experimental Results

The following section is a summary of some results acquired from literature to aid in choosing the model and approach to develop this project.

To emphasize the importance of the 4DRT-based perception module the author of this paper, (Paperswithcode.com, 2022) introduced a basic neural network for 3D object detection (referred to as the baseline NN) that takes 4DRT as its input. Their experiments on the K-RADAR dataset reveal that the 4DRT-based baseline NN excels in the task of 3D object detection, particularly in challenging weather conditions, outperforming the LiDAR-based network.

In another study carried out by (Scheiner et al., 2021), the objective is to perform object detection on automotive RADAR point clouds to identify moving road users using two end-to-end object detectors (YOLOv3 and PointPillars). YOLOv3 performs the best with a mean Average Precision (mAP) of 53.96%, offering potential for combining static and dynamic object detection in the future.

PointPillars, a point-cloud-based object detector, falls behind with a 36.89% mAP. Improving its performance to 45.82% by using an enhanced CNN backbone still falls short of YOLOv3-like results. While point-cloud-based detectors have potential, they are not yet on par with image-based variants. However, ongoing model development may bridge this performance gap, offering increased flexibility for various sensor types and improved speed in the future.

Another study comparing SSD and Faster R-CNN on RADAR Imagery done by (Stroescu et al., 2021), Concluded the training loss for SSD consistently decreases with each training epoch, the Faster R-CNN loss exhibits significantly lower values, leading to higher mean Average Precisions (mAPs). One plausible reason for the superior performance of Faster R-CNN can be learned from a study done by (Huang et al., 2016), which delves into the trade-off between speed and accuracy among various detectors on the COCO dataset. This research found that Faster R-CNN excels in accuracy, whereas SSD, though faster, struggles significantly when it comes to detecting small objects.

2.11 Conclusion

In progress...

Chapter 3 Analysis and Design

3.1 Application Overview

Object detection from 4D RADAR point clouds is a very challenging task for a final year thesis and it presents a massive set of unique challenges. These challenges stem from the nature of RADAR data, the complexity of the PV-RCNN architecture and the lack of industry knowledge in such a specific niche of software development for an undergraduate. This analysis aims to highlight these challenges to minimise the risk of running into too many blockers during the implementation phase.

3.2 Potential Blockers and Considerations

3.2.1 Data Sparsity and Quality

RADAR point clouds are inherently sparser than those obtained from LiDAR, which can make it difficult to detect and classify objects with high accuracy. Furthermore, RADAR data can be affected by noise and clutter from the environment, such as reflections from ground surfaces or nearby objects, making it challenging to isolate and identify relevant features for object detection.

3.2.2 Dataset Acquisition

Accessing datasets with 4D RADAR data and ground truth annotations, essential for developing and testing object detection models like PV-RCNN, presents a significant challenge. The limited availability of such datasets is often due to their proprietary nature and the competitive advantage they confer within the automotive and tech sectors. Companies heavily invest in sensor technology and data collection, treating their datasets as key assets to maintain competitiveness. Consequently, the protective stance on data sharing impedes collaborative research and development, especially for smaller entities and academics, slowing progress in object detection and autonomous vehicle technologies.

3.2.3 Computational Complexity

The advanced object detection models tailored for 4D RADAR point clouds demand high-end computing power, often requiring multiple GPUs due to their computational complexity, which involves voxelization, feature extraction, and integrating point and

voxel-level data. This need for substantial computational resources can result in prolonged training periods and slow development, potentially limiting the project's scale and scope. The iterative nature of machine learning, with its need for numerous experiments and parameter adjustments for optimal performance, exacerbates this challenge, especially for those without access to the computing capabilities of well-funded organizations.

3.2.4 Steep Learning Curve

The challenge of navigating a niche and technically complex task of object detection using 4D RADAR point clouds is particularly daunting for an undergraduate student, largely due to the steep learning curve and the lack of accessible industry knowledge. This specialized domain not only requires a deep understanding of advanced concepts in RADAR technology, computer vision, and machine/deep learning but also necessitates insights into the current state of the art and industry practices that are often not available in academic curricula or publicly accessible resources. The proprietary nature of much of the research and development in this area further exacerbates the difficulty in obtaining relevant and up-to-date information. This means that beyond mastering the technical skills, there's an additional layer of challenge in simply understanding the context, applications, and potential limitations of this project.

3.3 Scope

The project's scope for tackling the task of object detection, tracking and classification will need to be deliberately concentrated on the detection of objects within point clouds, due to the challenges explained in [\[3.2\]](#) and due to time constraints. This strategic limitation is critical to manage complexity, as it allows for a more manageable development process and provides a strong basis for future expansion into tracking and classification once detection is mastered.

3.4 Prerequisites

The following courses were done to gain a base knowledge for tackling this project:

1. Udemy. (2024). *Advanced Driver Assistance Systems (ADAS)*. [online] Available at: <https://www.udemy.com/course/advanced-driver-assistance-systems/learn/lecture/24251478?start=0#overview>
2. Udemy. (2024). *Automotive Radar*. [online] Available at: <https://www.udemy.com/course/automotive-radar-basics-to-advance/learn/lecture/20249534?start=15#overview>
3. Udemy. (2024). *Python for Computer Vision with OpenCV and Deep Learning*. [online] Available at: <https://www.udemy.com/course/python-for-computer-vision-with-opencv-and-deep-learning/learn/lecture/12257624?start=0#overview>
4. Udemy. (2024). *Deep Learning: Recurrent Neural Networks in Python*. [online] Available at: <https://www.udemy.com/course/deep-learning-recurrent-neural-networks-in-python/learn/lecture/21514852?start=0#overview>

3.5 Software Development Process

3.5.1 Agile

The author was grounded in Agile methodology. This approach was particularly suitable given the complexity and the exploratory nature of designing such an advanced neural network. Agile's iterative development cycles, known as sprints, allowed the author to break down the immense task into manageable segments, each delivering progress in increments. Embracing Agile's adaptive planning and development also provided the flexibility to respond to changes, which is crucial in a cutting-edge field where new discoveries and knowledge alter the course of development. This methodology, with its emphasis on continuous learning and adjustment, was essential in navigating the challenges of this project.

Key aspects of Agile methodology suited to this project are:

- An iterative and incremental approach to project management
- Emphasis on adaptability and flexibility
- Breaking down projects into smaller, manageable tasks
- Setting short-term goals and adjusting as needed
- Incorporating regular feedback and open communication.

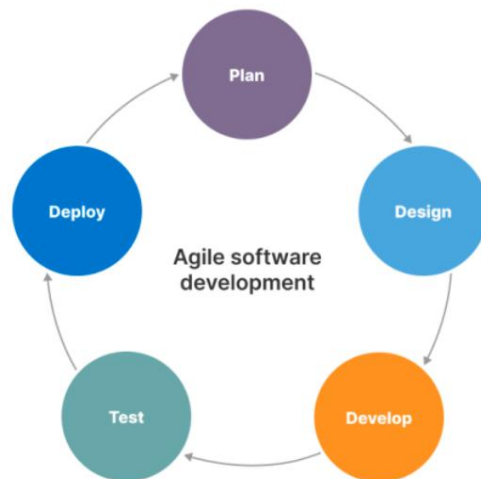


Figure 12: Agile Diagram (Aha, 2024)

3.5.2 Minimal viable product

Illustrated by **Figure 13** A Minimal Viable Product (MVP) is essentially a streamlined version of a product, designed to be functional enough for an initial release. It includes only the most essential features that define the product's core purpose, aiming to meet the needs of the first set of users, often referred to as early adopters. The primary objective of an MVP is to launch quickly and cost-effectively, allowing the development team to collect and analyse user feedback at an early stage. This feedback is invaluable as it guides subsequent development, ensuring that further enhancements and features are aligned with actual user needs and preferences. By focusing on the essentials, an MVP helps in validating product-market fit, minimizing initial investment, and reducing the risks associated with product development.

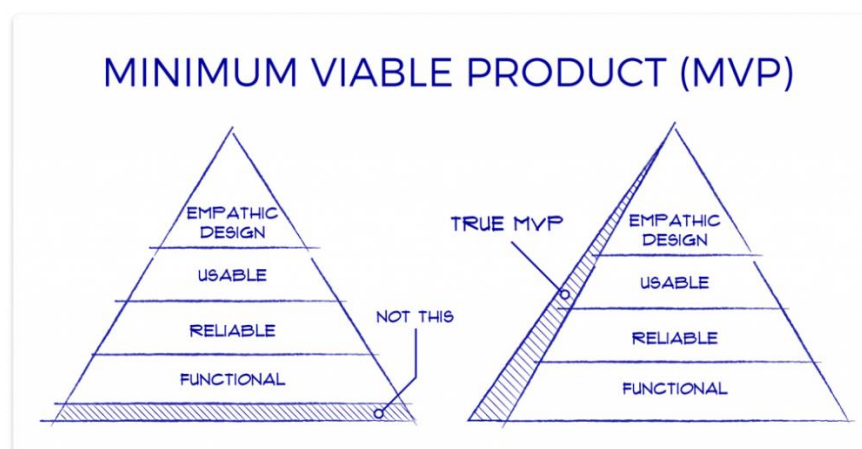


Figure 13: MVP Diagram

3.5.3 Project Management

The author leveraged ClickUp to implement agile methodology and manage the project by breaking down the workload into tasks and subtasks with urgency levels, due dates and assigning the tasks to sprints over the development process. ClickUp is a versatile project management and productivity tool that enables teams or solo developers to organize tasks, collaborate on projects, and track progress in a unified platform, facilitating the use of agile for efficient project development. ClickUp is a free alternative tool to Jira, which is the software of choice used by the company the author did his placement with.

3.6 Tools and Framework Considered

3.6.1 Python

What is Python?

Python is a powerful, high-level programming language that aims to be easy to read and understand. The term “high-level” refers to how closely related to human languages it is in comparison to other computer languages. Since the Python community is vibrant and active, learning this language will provide developers access to a wealth of useful tools that Python users have created since its inception.

Why Python?

Python's appeal to developers stems from its versatility across numerous fields like data science, AI, web and desktop app development, statistics, math, and scientific studies. The abundance of open-source libraries and the continuous growth of easy-to-use tools contribute to its popularity. The language is known for being beginner-friendly, which helps its already large development community grow even faster. Moreover, the vast Python community ensures that support, advice, and solutions are readily available for anyone encountering challenges with their projects.

What does Python's simplicity mean?

The best way to highlight Python's simplicity and user-friendliness is by contrasting how it and other programming languages, such as Java and C++, tackle the same basic task.

Consider the quintessential programming task of displaying "Hello World!" in the terminal as an example.

C++	Java	Python
<pre> 1. #include <iostream> 2. int main() 3. { 4. std::cout<<"Hello,world!\n"; 5. return 0; 6. }</pre>	<pre> 1. class HelloWorldApp { 2. public static void main(String[] args) { 3. System.out.println("Hello World!"); 4. } 5. }</pre>	<pre> 1. print("Hello World") 2.</pre>

3.6.2 Jupyter Notebook

Jupyter, short for Julia, Python, and R, began with these languages but now supports many more. The Jupyter Notebook is a free, open-source web application that allows for the sharing and collaborative editing of programming work. It enables programmers to create "notebooks," interactive documents that blend code, commentary, multimedia, and visuals, facilitating code execution directly in a web browser and proving valuable for educational demonstrations.

Advantages of using Jupyter Notebook:

- **Data visualisation.** Jupyter is often first encountered through its ability to display data sets as graphics, allowing for the creation, sharing, and dynamic modification of visualizations.
- **Code Share.** Cloud platforms like GitHub and Pastebin allow developers to exchange code but are generally inactive. Jupyter Notebook will enable developers to preview your code, run it, and examine the results immediately in your browser.
- **Live interaction with code.** The code in Jupyter Notebook is not static; it is real-time, gradually editable, and replay-able, with feedback provided immediately in the browser. Notebooks can have user controls that can be utilised as code input sources.
- **Documenting code samples.** A developer might put some code in a Jupyter Notebook to show how it works step by step with real-time feedback. The code is

still fully functional, and the developer can add interaction by explaining, delivering, and talking simultaneously.

A Jupyter Notebook can have many parts, and each one is made up of different blocks.

Components of Jupyter Notebook:

- **Text and HTML** Anywhere on the page, developers can enter plain text or content written in Markdown syntax to turn it into HTML. The notebook template can have a built-in CSS style or be added to it.
- **Code and output.** Although developers can add support for other languages in the Jupyter environment, such as R or Julia, Jupyter Notebooks' programming is usually written in Python. The code blocks can be run and repeated in whatever order as many times as desired, and the results of the executed code appear immediately after the code blocks.
- **Multimedia** Due to the fact that it is built on web technologies, Jupyter Notebook can display supported forms of multimedia on a web page. They can be pre-programmed by developers with the help of the IPython.display module, or they can be inserted into a notebook as HTML elements.
- **Data** In addition to the `.ipynb` file that makes up a Jupyter Notebook, data can be given as a separate file or imported programmatically. For instance, code might be inserted into the notebook to download data from a public Internet repository or access a database connection.

The restriction set by Jupyter Notebook

- **Notebooks are not self-contained.** The fact that Jupyter Notebook requires the Jupyter runtime and the libraries the developer wants to use is its biggest drawback. There are a few ways to create independent Jupyter Notebooks, but none of them is supported by the project. It's best to install or give laptops to those with the necessary infrastructure (via Anaconda, for example).
- **The session state is difficult to save.** Using the Jupyter Notebook toolset, you can't save the current state of the code in a Jupyter notebook and then load it again. Every time you load the notebook, the developer must run the code again to get it back to the way it was.

- **There is no interactive debugging or other IDE functionality.** Jupyter Notebook does not provide a full Python programming environment. Many functions that users anticipate from an IDE, such as interactive debugging, code completion, and module management, are missing.

3.7 Dataset

Several datasets with 4D radar have been released and applied in recent years. However, despite the lack of publicly available datasets, an analysis was conducted on the datasets that were available. From the analysis the author created the following summary:

Astyx, an early-released dataset, provides rich data for 3D object detection but is limited by its small size of 546 frames and 3000 object annotations, lacking special scenarios and urban data. RADIAL offers a medium-scale dataset with urban streets and highways but lacks 3D bounding boxes and tracking IDs and does not cover adverse weather conditions. View-of-Delft addresses the object tracking problems present in the other datasets with 8,693 frames and 120,000 annotated objects but has a short detection range and lacks 4D radar information for long-range mode. TJ4DRaDSet includes various driving scenarios but lacks data in middle and short-range modes and scenarios with adverse weather conditions. K-Radar provides rich driving scenarios with adverse weather conditions but lacks 4D radar point cloud in long-range mode.

View-of-delft was chosen for this experiment setup as it “consists of more than 123000 3D bounding box annotations, including more than 26000 pedestrian, 10000 cyclist and 26000 car labels.” (GitHub, 2024) Which is great for an object detection problem.

The View-of-delft researchers have an open-source GitHub repository containing useful guides to their dataset along with visualisation and evaluation scripts they used when conducting their research, this repository would come in very useful when exploring the dataset and getting familiar with the process of doing object detection with 4D radar point clouds.

3.7.1 View-Of-Delft Sensor Setup

The radar sensor utilized is a ZF FRGen21 3+1D radar, operating at approximately 13 Hz, and is mounted behind the vehicle's front bumper. The provided radar point clouds undergo ego-motion compensation to account for any motion between radar and camera data capture, ensuring consistency when overlaying both datasets. (GitHub, 2024)



Figure 12: Prius Sensor Setup VOD (GitHub, 2024)

3.7.2 Dataset Frame Information

These radar point clouds are stored in bin files, with each file containing a set of points represented as a $N \times 7$ array, where N is the number of points, and 7 denotes the number of features: $[x, y, z, RCS, v_r, v_{r_compensated}, time]$. Here, " v_r " signifies the relative radial velocity, " $v_{r_compensated}$ " indicates the absolute radial velocity compensated for ego motion, and " $time$ " denotes the point's time ID, indicating its originating scan (GitHub, 2024).

3.7.3 Dataset Label Information

Sensor fusion in autonomous systems integrates data from multiple sensors like LiDAR and radar to enhance accuracy and robustness, utilizing LiDAR's high-resolution for precise object labelling which is applied to radar data for additional insights such as velocity. Aligning sensor data in both space and time is crucial, often using LiDAR as the

benchmark for its detailed spatial data. LiDAR's clear 3D point clouds simplify labelling tasks, with labels easily transferred to radar data for consistency. The dataset's structure prioritizes LiDAR, reflecting its initial design focus and convenience in labelling, leading to radar data being adjunct and reliant on LiDAR for annotations.

For frame number 1047, the first line of the label information could be interpreted as follows:

- **Object Type:** rider.
- **Truncation:** 1 (possibly indicating high confidence or manual verification)
- **Occlusion:** 0 (fully visible)
- **Observation Angle:** 1.716500830699201
- **Bounding Box:** 979.41486 789.5281 1018.06165 866.89154
- **3D Object Dimensions:** 1.503325462332693 (height), 0.7167884312694952 (width), 0.6358283468841199 (length)
- **3D Object Location:** 0.7805723338707173 (x), 4.960184749066411 (y), 31.026849236059597 (z)
- **Rotation Y:** -4.541531818868102
- **Score:** 1 (high confidence or manually verified)

3.8 Data Visualization and Exploration

Before any work can be carried out, the radar scans need to be visualised in a point cloud for easier interpretation and comparing this visualisation to the output of the camera. K3D was used by VOD, however the author chose Open3D because it is preferred for its extensive feature set, active community support, cross-platform compatibility, seamless Python integration, and ongoing development. These factors make it a versatile choice for various 3D data processing tasks compared to K3D. Once the point cloud is visualised its good practice to get the annotations from the dataset and attempt to draw the 3D bounding boxes inside of the point cloud to be able to interpret the different objects (car, pedestrian, cyclist etc.) present in the scene.

It would also be beneficial to calculate the number of points within a specific scene and any output any other relevant information.

3.9 Model Architecture Considered

Following an in-depth comparison of object detection architectures for 4D radar based on the literature available and comparing results, it was decided to use the PV-RCNN architecture from the OpenPCDet library as a guide and reference.

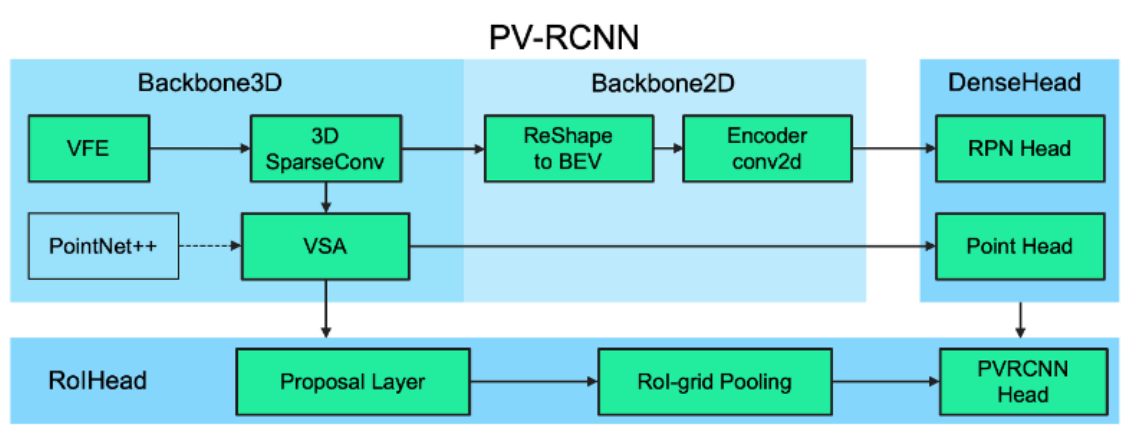


Fig 14: OpenPCDet PV-RCNN Architecture Diagram (GitHub, 2024)

The PV-RCNN architecture combines 3D voxel CNNs and PointNet for precise 3D object detection from point clouds. It leverages voxel CNNs for efficient feature learning and high-quality proposal generation, while incorporating Point Net’s ability to capture detailed contextual information. The framework introduces voxel-to-key point scene encoding to condense scene features into key points and point-to-grid RoI feature abstraction for refining proposal confidence and location, effectively integrating the strengths of both network types for enhanced performance.

3.9.1 PV-RCNN Components Overview

In this section the author breaks down the architecture into the main components based on **Figure 13**, to provide a simple understanding for the reader.

1. **VFE (Voxel Feature Encoding):** This is the initial stage where raw point clouds are converted into a structured voxel representation. The VFE learns to encode the raw point cloud data into a fixed-size feature representation for each voxel, which can capture the essential information while reducing the sparsity of point clouds.
2. **Backbone3D**
3D SparseConv: Once the point clouds are voxelized, the 3D Sparse Convolutional Network processes these voxels to generate high-dimensional

sparse feature volumes. It is designed to efficiently handle the sparsity in voxelized point clouds by only computing features at occupied voxels.

VSA (Voxel Set Abstraction): The VSA module is responsible for summarizing the voxelized features into a smaller set of key points. This abstraction process leverages the strength of PointNet++ to learn a more discriminative feature representation from the unordered point set within each voxel.

3. **Backbone2D**

Reshape to BEV (Bird's Eye View): The high-dimensional features from the 3D backbone are reshaped into a 2D representation corresponding to the bird's eye view of the scene. This process essentially projects the 3D features onto a 2D plane.

Encoder conv2d: This 2D feature map is then processed by a 2D convolutional encoder which further refines the features and captures spatial relationships in the bird's eye view.

4. **DenseHead**

RPN Head (Region Proposal Network): The RPN Head uses the refined 2D feature maps to generate region proposals. These proposals are candidate regions where objects might be present.

Point Head: Parallel to the RPN Head, the Point Head processes the point features from the VSA to generate key point features that describe the objects in the scene.

5. **RoI Head**

Proposal Layer: This layer takes the region proposals from the RPN Head and the key point features from the Point Head to generate accurate 3D bounding box proposals.

RoI-grid Pooling: The RoI-grid Pooling module aggregates the key point features onto a structured grid within each 3D RoI using set abstraction operations. This step is crucial for capturing the local context around each proposal.

6. **PVRCNN Head:** Finally, the PV-RCNN Head combines the features from the RoI-grid Pooling module with the proposals from the RoI Head to refine the proposals and predict the final bounding boxes along with the object classification.

3.10 Training and Inference Details

The PV-RCNN framework from OpenPCDet undergoes training from the ground up in a seamless end-to-end process utilizing the ADAM optimization algorithm. The researchers utilized the KITTI dataset to train the entire network using a batch size of 24 and a learning rate of 0.01 across 80 epochs. This training is performed on 8 GTX 1080 Ti graphics processing units and completes in approximately 5 hours.

Given that the author of this thesis, training setup consists of a laptop with a single Nvidia 1660 Ti GPU, the author may need to consider downsizing the dataset and be prepared for the training duration to extend over several days. This is due to the lower computational power compared to the original setup with 8 GTX 1080 Ti GPUs. Alternatively, to maintain the scale of the dataset and potentially accelerate the training process, the author could leverage a cloud computing service that offers more powerful GPU resources. This would allow me to conduct the training more efficiently, albeit at a potential increase in cost. The author will also need to experiment with hyperparameters such as the epochs.

3.11 Evaluation Metrics

Based on research the author chose the mean Average Precision (mAP), (illustrated by **Figure 15**) metric for evaluating the effectiveness of the model, it utilizes the following subsidiary metrics:

1. Confusion Matrix
2. Intersection over Union (IoU)
3. Precision
4. Recall

It is the average of the Average Precision (AP) for each class. AP for a single class is calculated by plotting a Precision-Recall (P-R) curve, which graphs the precision of the model against its recall at different thresholds. Precision measures the accuracy of the positive predictions, while recall quantifies the model's ability to identify all relevant instances. For each class, AP is computed as the area under the P-R curve, reflecting the model's accuracy across all levels of recall. The mAP is then determined by taking the mean of the AP scores across all classes. This provides a single performance summary

that balances both the precision and the ability of the model to detect each class, making it a comprehensive metric for tasks with multiple classes to be detected.

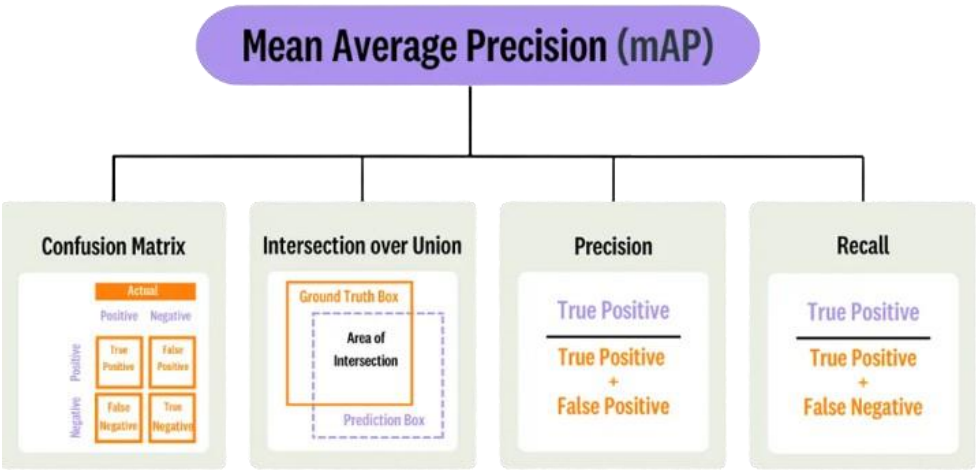


Fig 15: mAP Overview (kili-website, 2023)

mAP, or mean Average Precision, is often preferred over other evaluation metrics for object detection tasks due to its comprehensive assessment of both precision and recall across various levels of confidence thresholds. Unlike single-point metrics like accuracy, mAP considers the entire precision-recall curve, providing a more nuanced understanding of a model's performance. This makes mAP particularly suitable for scenarios where a balance between precision and recall is crucial, such as in real-world applications where false positives and false negatives carry different consequences. Additionally, mAP enables fair comparisons between different models by accounting for their performance across all possible operating points, offering a robust measure of overall detection quality.

Chapter 4 Implementation

4.1 Introduction

The chapter's primary objective is to give the reader a thorough grasp of the project's implementation procedure and the technologies employed. This chapter will cover the hardware and software utilised and the factors in its selection. It describes the project's setup and highlights the components of technical interest to the project, along with the challenges faced and how they were addressed.

4.2 Tools Used

This section outlines the authors choice of software tools, frameworks, and libraries used to carry out the research and project implementation.

4.2.1 Integrated Development Environment (IDEs)

The author chose PyCharm as is it is often the preferred IDE for implementing complex projects like implanting a PV-RCNN architecture for 4D radar due to its comprehensive Python-centric features. It offers an integrated environment tailored for Python development, including smart code navigation, advanced debugging, and refactoring tools, which significantly enhance productivity and code quality. PyCharm's support for scientific libraries and frameworks, such as TensorFlow and PyTorch, streamlines the machine learning workflow, from data exploration to model training and evaluation. The IDE's virtual environment management simplifies dependency handling, ensuring project consistency. Moreover, PyCharm's powerful visualization capabilities, coupled with its interactive Python console, facilitate the examination and interpretation of 4D radar data, making it an invaluable tool for researchers and developers working on advanced radar object detection systems.

4.2.2 Ubuntu

After an initial attempt was made to setup the OpenPCDet repository on a windows system the author found an external resource that stated this repository is only supported on Linux, despite there being no mention of this in their repository and a lot of lost time the author needed to setup Linux on an Ubuntu Virtual Machine.

OpenPCDet, a library tailored for point cloud detection tasks in autonomous driving, operates on Linux due to its reliance on Linux-specific libraries, tools, and a development

environment that's closely aligned with the needs of high-performance computing tasks, such as those involving CUDA for GPU acceleration.

4.2.3 Libraries

An outline of the main libraries needed:

Previously mentioned in [\[3.7\]](#) **Open3d** was chosen for visualising the point cloud.

Numpy: A fundamental package for scientific computing with Python, providing support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on these arrays.

Numba: An open-source JIT compiler that translates a subset of Python and NumPy code into fast machine code, significantly accelerating execution.

PyTorch (Version 1.1): A deep learning library that provides a flexible and powerful array library, Tensor, with GPU acceleration and automatic differentiation capabilities for building and training neural networks.

TensorboardX: An extension to TensorBoard, providing visualization and tooling needed for machine learning experimentation, such as tracking and visualizing metrics, and model graphs.

PyYAML: A YAML parser and emitter for Python, enabling easy reading and writing of YAML files for configuration or data serialization.

Scikit-image: A collection of algorithms for image processing in Python, providing tools for image manipulation and analysis.

Torchvision: A package consisting of popular datasets, model architectures, and common image transformations for computer vision.

OpenCV-Python: A Python wrapper for OpenCV, offering access to a wide range of image processing and computer vision functions.

4.3 Source Control

Git is the source control technology used, and the source is maintained in a GitHub repository @ <https://github.com/nicktmv/four-d-radar-thesis>

4.4 Installation

To install and run the project.

1. Clone the repository.

```
git clone https://github.com/nicktmv/four-d-radar-thesis.git
```

2. Install the Python dependencies.

```
pip install -r requirements.txt
```

3. Rest pending.

4.5 Data Loader & Split

To facilitate model training and validation on the dataset which is 26.1 GB in total, it is partitioned into training and testing subsets. This partitioning is guided by predefined indices listed in 'train.txt' and 'test.txt' files, ensuring a consistent and reproducible split across different experimental runs. This approach adheres to best practices in machine learning by preventing data leakage and ensuring the model's performance is evaluated on unseen data.

Data loader code...


```

Run: data_loader x
D:\FYP-K00200158\venv\Scripts\python.exe D:\FYP-K00200158\venv\Scripts\python.exe
Loading training data...
Loaded 16 labels for 00544.bin
Loaded 16 labels for 00545.bin
Loaded 16 labels for 00546.bin
Loaded 16 labels for 00547.bin
Loaded 16 labels for 00548.bin
Loaded 15 labels for 00549.bin
Loaded 15 labels for 00550.bin
Loaded 15 labels for 00551.bin
Loaded 15 labels for 00552.bin
Loaded 15 labels for 00553.bin
Loaded 15 labels for 00554.bin
Loaded 13 labels for 00555.bin
Loaded 13 labels for 00556.bin
Loaded 13 labels for 00557.bin
Loaded 13 labels for 00558.bin
Loaded 13 labels for 00559.bin
Loaded 13 labels for 00560.bin
Loaded 12 labels for 00561.bin
Loaded 11 labels for 00562.bin
Loaded 11 labels for 00563.bin
Loaded 11 labels for 00564.bin
Loaded 10 labels for 00565.bin
Loaded 10 labels for 00566.bin
Loaded 9 labels for 00567.bin
Loaded 9 labels for 00568.bin

Run: data_loader x
Loaded 1 labels for 09912.bin
Loaded 1 labels for 09913.bin
Loaded 1 labels for 09914.bin
Loaded 1 labels for 09915.bin
Loaded 1 labels for 09916.bin
Loaded 1 labels for 09917.bin
Loaded 1 labels for 09918.bin
Loaded 1 labels for 09919.bin
Loaded 1 labels for 09920.bin
Loaded 1 labels for 09921.bin
Loaded 1 labels for 09922.bin
Loaded 1 labels for 09923.bin
Loaded 1 labels for 09924.bin
Loaded 1 labels for 09925.bin
Loaded 1 labels for 09926.bin
Loaded 1 labels for 09927.bin
Loaded 2 labels for 09928.bin
Loaded 2 labels for 09929.bin
Loaded 2 labels for 09930.bin
Labels were loaded for 5139 out of 5139 files.
Number of points in the first training file: 320

Loading testing data...
Number of points in the first testing file: 349

Process finished with exit code 0

```

Fig 16: Data Loader output

Illustrated by **Fig 16:** An output for the number of labels loaded for each .bin file in the training set was given. Then the loader gets a total number of the labels loaded for the training set, along with the number of points in the first point cloud. Lastly the loader attempts to load in the testing set and prints out the number of the points in the first .bin file.

During the data loading process and assumption is made “Assuming 7 features (which is correct since the point cloud is a $N7$ array) \times 4 bytes each”

This assumption is verified in check_data_structure.py by the following code:

```

1. def main():
2.     file_path = '../data/view_of_delft_PUBLIC/radar/training/velodyne/00544.bin'
3.
4.     # Ensure the file exists
5.     if not os.path.exists(file_path):
6.         print(f"File not found: {file_path}")
7.         return
8.
9.     try:
10.        with open(file_path, 'rb') as file:
11.            data = file.read() # Read the entire file
12.            # Calculate the number of points assuming each point is represented by 28
13.            bytes (7 features x 4 bytes each)
14.            num_points = len(data) // 28
15.            remainder = len(data) % 28
16.            # Print results
17.            print(f"Total data length in bytes: {len(data)}")
18.            print(f"Number of points calculated: {num_points}")
19.            print(f"Remainder (should be 0 if assumption is correct): {remainder}")
20.
21.        # Check a few points to see if the assumption holds

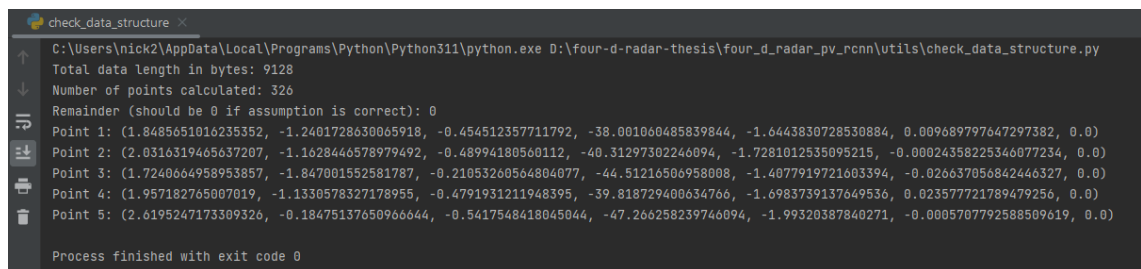
```

```

21.         for i in range(min(num_points, 5)): # Check the first 5 points or the total
number of points if less than 5
22.             point_data = data[i * 28:(i + 1) * 28] # Extract the data for one point
23.             point = struct.unpack('7f', point_data) # Unpack the point data
24.             print(f"Point {i + 1}: {point}")
25.
26.     except Exception as e:
27.         print(f"An error occurred: {e}")

```

Verified by using a frame 00544, the assumption is true as the remainder should ideally be 0 if the assumption that each point is exactly 28 bytes, which is correct, illustrated by **Figure 17**



```

check_data_structure
C:\Users\nick2\AppData\Local\Programs\Python\Python311\python.exe D:\four-d-radar-thesis\four_d_radar_pv_rcnn\utils\check_data_structure.py
Total data length in bytes: 9128
Number of points calculated: 326
Remainder (should be 0 if assumption is correct): 0
Point 1: (1.8485651016235352, -1.2401728630065918, -0.454512357711792, -38.001060485839844, -1.6443830728530884, 0.009689797647297382, 0.0)
Point 2: (2.0316319465637207, -1.1628446578979492, -0.48994180560112, -40.31297302246094, -1.7281012535095215, -0.00024358225346077234, 0.0)
Point 3: (1.7240664958953857, -1.847001552581787, -0.21053260564804077, -44.51216506958008, -1.4077919721603394, -0.026637056842446327, 0.0)
Point 4: (1.957182765007019, -1.1330578327178955, -0.4791931211948395, -39.818729400634766, -1.6983739137649536, 0.023577721789479256, 0.0)
Point 5: (2.6195247173309326, -0.18475137650966644, -0.5417548418045044, -47.266258239746094, -1.99320387840271, -0.0005707792588509619, 0.0)
Process finished with exit code 0

```

Fig 17: check_data_structure.py output

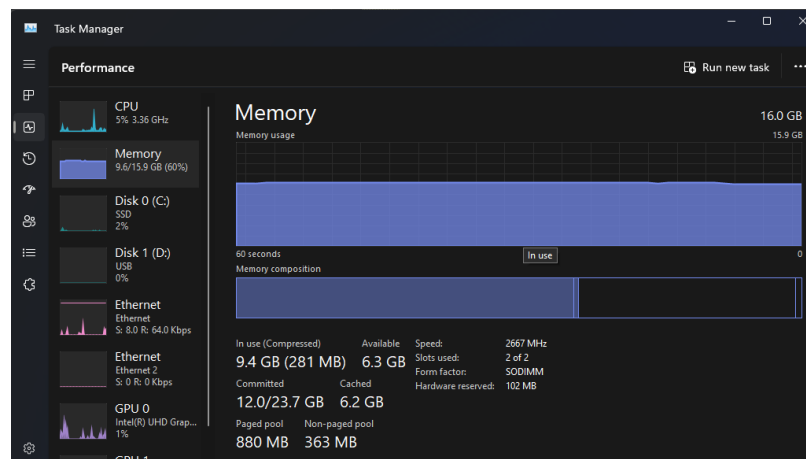


Fig 17: Task Manager screenshot

Loading and visualizing the large-scale point cloud data consumes substantial memory, as evident from **Fig 17:** screenshot showing 60% usage of the available 16 GB RAM. This is due to the storage needs of the extensive data and its labels in RAM, necessary for real-time processing and visualization. The considerable memory footprint is influenced by the data's volume and complexity, along with the memory overhead of the utilized data structures.

4.6 Point Cloud Visualization in Open3d




Fig 14: Visualization of radar point cloud for Frame 01047

label information is being interpreted incorrectly and used to create the bounding boxes in the 3D visualization.




Fig 14: Visualization of radar point cloud for Frame 01047



Fig 15: Camera output for Frame 01047

```
1. (ratrack) nick@nick-VirtualBox:~/Documents/GitHub/RaTrack$
/home/nick/Documents/GitHub/RaTrack/lib/setup.py --user bash:
/home/nick/Documents/GitHub/RaTrack/lib/setup.py: Permission denied
2.
```

```
Windows PowerShell
exit code: 1
[12 lines of output]
running bdist_wheel
running build
running build_ext
building 'SharedArray' extension
creating build
creating build\temp.win-amd64-cpython-37
creating build\temp.win-amd64-cpython-37\Release
creating build\temp.win-amd64-cpython-37\Release\src
"C:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.37.32822\bin\HostX86\x64\cl.exe" /c /nologo /O2 /W3 /GL /DNDEBUG /MD -IC:\Users\nick2\AppData\Local\Temp\pip-build-env-ricdb5h9\overlay\Lib\site-packages\numpy\core\include -IC:\ProgramData\Anaconda3\include -IC:\ProgramData\Anaconda3\include -IC:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.37.32822\include" "-IC:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.37.32822\ATLMFC\include" "-IC:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.37.32822\ATLMFC\include" "-IC:\Program Files\Microsoft Visual Studio\2022\Community\VC\Auxiliary\VS\include" "-IC:\Program Files (x86)\Windows Kits\10\include\10.0.22621.0\ucrt" "-IC:\Program Files (x86)\Windows Kits\10\include\10.0.22621.0\um" "-IC:\Program Files (x86)\Windows Kits\10\include\10.0.22621.0\shared" "-IC:\Program Files (x86)\Windows Kits\10\include\10.0.22621.0\winrt" "-IC:\Program Files (x86)\Windows Kits\10\include\10.0.22621.0\cppwinrt" /Tc.\src\map_owner.c /Fobuild\temp.win-amd64-cpython-37\Release\.\src\map_owner.obj
map_owner.c
.\src\map_owner.c(19): fatal error C1083: Cannot open include file: 'sys/mman.h': No such file or directory
error: command 'C:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.37.32822\bin\HostX86\x64\cl.exe' failed with exit code 2
[end of output]

note: This error originates from a subprocess, and is likely not a problem with pip.
ERROR: Failed building wheel for SharedArray
Failed to build SharedArray
ERROR: Could not build wheels for SharedArray, which is required to install pyproject.toml-based projects
PS D:\OpenPCDet> SS
```

- 4.7 Getting required repositories
- 4.8 GitHub Contributions
- 4.9 Building the model
- 4.10 Training the model

Chapter 5 Testing and Results

5.1 Introduction

Unit testing was utilised to isolate and test individual code components for this application during the testing process. This approach allowed for the verification of the application's functionality in isolation, making it easier to identify and address any issues that arose.

5.2 Unit Testing

5.2.1 What is Unit Testing, and why is important?

Unit testing is a software development technique that involves testing individual components or units of a software application to verify their functionality. In this approach, each unit is tested in isolation, with dependencies typically mocked or stubbed to focus solely on the unit's behaviour. The significance of unit testing lies in its ability to detect errors early in the development cycle, reducing the complexity and cost associated with later-stage bug fixes. By isolating and testing units separately, developers can swiftly identify and resolve issues that might be elusive in broader testing strategies. Furthermore, unit testing plays a crucial role in maintaining code integrity during updates, preventing new modifications from disrupting existing functionality. Integrating unit tests into a continuous integration workflow ensures thorough vetting of changes prior to merging them into the main code repository. Additionally, unit testing promotes the development of clean, maintainable code by encouraging modular design, which simplifies understanding and maintenance. Ultimately, unit testing is a vital practice in software development, enhancing the quality, stability, and reliability of software products, while fostering best coding practices and preventing regression.

5.2.2 Testing Framework in PyCharm

The author utilized PyCharm's Python test framework to streamline the testing process, leveraging its integrated tools to efficiently write, manage, and execute unit tests directly within the IDE, ensuring each component functioned correctly before integration.

5.3 Unit Testing Done

Pending...

5.4 Result Analysis

No results...

5.5 Discussion of Findings

Pending...

Chapter 6 Conclusion:

Pending...

6.1 Future Work

Pending

Chapter 7 Appendices

Zhou, Y., Liu, L., Zhao, H., López-Benítez, M., Yu, L. and Yue, Y. (2022). Towards Deep Radar Perception for Autonomous Driving: Datasets, Methods, and Challenges. *Sensors*, [online] 22(11), pp.4208–4208. doi:<https://doi.org/10.3390/s22114208>.

2.0 A Vision for Safety AUTOMATED DRIVING SYSTEMS. (n.d.). Available at: https://www.nhtsa.gov/sites/nhtsa.gov/files/documents/13069a-ads2.0_090617_v9a_tag.pdf.

Anwesh Marwade (2020). *Kalman Filtering: An Intuitive Guide Based on Bayesian Approach*. [online] Medium. Available at: <https://towardsdatascience.com/kalman-filtering-an-intuitive-guide-based-on-bayesian-approach-49c78b843ac7> [Accessed 10 Oct. 2023].

AR, A. (2021). *LiDAR vs RADAR: Detection, Tracking, and Imaging - For 4D Sensing, AI, AR, AV... - Medium*. [online] Medium. Available at: <https://4sense.medium.com/lidar-vs-radar-detection-tracking-and-imaging-ca528c0e9aae> [Accessed 24 Oct. 2023].

AUTOCRYPT. (2023). *The State of Level 3 Autonomous Driving in 2023 / AUTOCRYPT*. [online] Available at: <https://autocrypt.io/the-state-of-level-3-autonomous-driving-in-2023/> [Accessed 3 Oct. 2023].

Avijeet Biswal (2020). *Power of Recurrent Neural Networks (RNN): Revolutionizing AI*. [online] Simplilearn.com. Available at: <https://www.simplilearn.com/tutorials/deep-learning-tutorial/rnn> [Accessed 1 Nov. 2023].

baeldung (2022). *The Viola-Jones Algorithm / Baeldung on Computer Science*. [online] Baeldung on Computer Science. Available at: <https://www.baeldung.com/cs/viola-jones-algorithm> [Accessed 25 Oct. 2023].

Balasubramaniam, A. and Pasricha, S. (2022). *Object Detection in Autonomous Vehicles: Status and Open Challenges*. [online] ResearchGate. Available at: https://www.researchgate.net/publication/357953408_Object_Detection_in_Autonomous_Vehicles_Status_and_Open_Challenges#pf2 [Accessed 25 Oct. 2023].

Balasubramaniam, A. and Pasricha, S. (n.d.). *Object Detection in Autonomous Vehicles: Status and Open Challenges*. [online] Available at:

<https://arxiv.org/ftp/arxiv/papers/2201/2201.07706.pdf#:~:text=Object%20detection%20consists%20of%20two>

Barnard, M. (2016). *Tesla & Google Disagree About LIDAR - Which Is Right?* - *CleanTechnica*. [online] CleanTechnica. Available at: <https://cleantechnica.com/2016/07/29/tesla-google-disagree-lidar-right/> [Accessed 10 Oct. 2023].

Bloom, C. (2020). *Introduction to Radar*. [online] Arrow.com. Available at: <https://www.arrow.com/en/research-and-events/articles/introduction-to-radar> [Accessed 23 Oct. 2023].

Borah, C. (2020). *Evolution of Object Detection - Analytics Vidhya - Medium*. [online] Medium. Available at: <https://medium.com/analytics-vidhya/evolution-of-object-detection-582259d2aa9b> [Accessed 25 Oct. 2023].

Cohen, J. (2020). *How RADARs work*. [online] Welcome to The Library! Available at: <https://www.thinkautonomous.ai/blog/how-radars-work/> [Accessed 4 Oct. 2023].

Cohen, J. (2023). *4D LiDARs vs 4D RADARs — Why the LiDAR vs RADAR comparison is more relevant today than ever*. [online] Welcome to The Library! Available at: <https://www.thinkautonomous.ai/blog/fmcw-lidars-vs-imaging-radars/#:~:text=4D%20RADARs%20work%20using%20MIMO,have%20a%20pretty%20bad%20resolution>. [Accessed 10 Oct. 2023].

Davies, E.R. (2022). The dramatically changing face of computer vision. *Elsevier eBooks*, [online] pp.1–91. doi:<https://doi.org/10.1016/b978-0-12-822109-9.00010-2>.

Doppler effect | Definition, Example, & Facts | Britannica. (2023). In: *Encyclopædia Britannica*. [online] Available at: <https://www.britannica.com/science/Doppler-effect> [Accessed 4 Oct. 2023].

Everythingrf.com. (2021). *What are 4D Radars? - everything RF*. [online] Available at: <https://www.everythingrf.com/community/what-are-4d-radars> [Accessed 10 Oct. 2023].

Gandhi, R. (2018). *R-CNN, Fast R-CNN, Faster R-CNN, YOLO — Object Detection Algorithms*. [online] Medium. Available at: <https://towardsdatascience.com/r-cnn-fast-r-cnn-faster-r-cnn-yolo-object-detection-algorithms-36d53571365e>

GeeksforGeeks. (2018). *Introduction to Recurrent Neural Network*. [online] Available at: <https://www.geeksforgeeks.org/introduction-to-recurrent-neural-network/> [Accessed 1 Nov. 2023].

Gibiansky, A. (2014). *Convolutional Neural Networks - Andrew Gibiansky*. [online] Gibiansky.com. Available at: [https://andrew.gibiansky.com/blog/machine-learning/convolutional-neural-networks/#:~:text=If%20we%20use%20an%20m,\(j%2Bb\).](https://andrew.gibiansky.com/blog/machine-learning/convolutional-neural-networks/#:~:text=If%20we%20use%20an%20m,(j%2Bb).)

Gillis, A.S., Burns, E. and Brush, K. (2023). *deep learning*. [online] Enterprise AI. Available at: <https://www.techtarget.com/searchenterpriseai/definition/deep-learning-deep-neural-network>

Girshick, R., Donahue, J., Darrell, T. and Malik, J. (2016). Region-Based Convolutional Networks for Accurate Object Detection and Segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(1), pp.142–158. doi:<https://doi.org/10.1109/tpami.2015.2437384>.

GitHub. (2023). *view-of-delft-dataset/figures/example_frame_2.png at main · tudelft-iv/view-of-delft-dataset*. [online] Available at: https://github.com/tudelft-iv/view-of-delft-dataset/blob/main/figures/example_frame_2.png [Accessed 24 Oct. 2023].

Hearst, M.A., Dumais, S.T., Osman, M., Platt, J. and Bernhard Schölkopf (1998). Support vector machines. *IEEE Intelligent Systems & Their Applications*, [online] 13(4), pp.18–28. doi:<https://doi.org/10.1109/5254.708428>.

Hesai Webmaster (2023). *What You Need to Know About Lidar: The Strengths and Limitations of Camera, Radar, and Lidar*. / HESAI. [online] HESAI. Available at: <https://www.hesaitech.com/what-you-need-to-know-about-lidar-the-strengths-and-limitations-of-camera-radar-and-lidar/#:~:text=Compared%20to%20the%20traditional%203D,over%20100%2C000%20points%20per%20frame.> [Accessed 24 Oct. 2023].

Jessica Van Brummelen, O'Brien, M., Gruyer, D. and Homayoun Najjaran (2018). Autonomous vehicle perception: The technology of today and tomorrow. *Transportation Research Part C-emerging Technologies*, [online] 89, pp.384–406. doi:<https://doi.org/10.1016/j.trc.2018.02.012>.

Levity.ai. (2023). *Deep Learning vs. Machine Learning – What's The Difference?* [online] Available at: <https://levity.ai/blog/difference-machine-learning-deep-learning>

Mathworks.com. (2023). *What Is SLAM (Simultaneous Localization and Mapping) – MATLAB & Simulink.* [online] Available at: [https://www.mathworks.com/discovery/slam.html#:~:text=SLAM%20\(simultaneous%20localization%20and%20mapping\)%20is%20a%20method%20used%20for,to%20map%20out%20unknown%20environments.](https://www.mathworks.com/discovery/slam.html#:~:text=SLAM%20(simultaneous%20localization%20and%20mapping)%20is%20a%20method%20used%20for,to%20map%20out%20unknown%20environments.) [Accessed 10 Oct. 2023].

Navtech Radar. (2023). *FMCW Radar.* [online] Available at: [https://navtechradar.com/explore/fmcw-radar/#:~:text=Frequency%20Modulated%20Continuous%20Wave%20\(FMCW,by%20the%20radar%20and%20compared.](https://navtechradar.com/explore/fmcw-radar/#:~:text=Frequency%20Modulated%20Continuous%20Wave%20(FMCW,by%20the%20radar%20and%20compared.) [Accessed 10 Oct. 2023].

NHTSA. (2020). *NHTSA.* [online] Available at: <https://www.nhtsa.gov/technology-innovation/automated-vehicles-safety> [Accessed 23 Oct. 2023].

Ouaknine, A. (2022). *Deep learning for radar data exploitation of autonomous vehicle.* [online] arXiv.org. Available at: <https://arxiv.org/abs/2203.08038>

ResearchGate. (2019). *Figure 1. Schematic diagram of a basic convolutional neural network...* [online] Available at: https://www.researchgate.net/figure/Schematic-diagram-of-a-basic-convolutional-neural-network-CNN-architecture-26_fig1_336805909

Synopsys.com. (2023). *What is an Autonomous Car? – How Self-Driving Cars Work / Synopsys.* [online] Available at: <https://www.synopsys.com/automotive/what-is-autonomous-car.html#:~:text=Definition,in%20the%20vehicle%20at%20all.> [Accessed 10 Oct. 2023].

Tyagi, M. (2021). *HOG (Histogram of Oriented Gradients): An Overview - Towards Data Science.* [online] Medium. Available at: <https://towardsdatascience.com/hog-histogram-of-oriented-gradients-67ecd887675f> [Accessed 25 Oct. 2023].

Udemy. (2023). *Automotive Radar.* [online] Available at: <https://www.udemy.com/course/automotive-radar-basics-to-advance/> [Accessed 10 Oct. 2023].

Xx, N., Xx, X. and Xxxx (n.d.). Millimeter Wave Sensing: A Review of Application Pipelines and Building Blocks. *IEEE SENSORS JOURNAL*, [online] (1). Available at: <https://arxiv.org/pdf/2012.13664.pdf>.

Yang, B., Guo, R., Liang, M., Casas, S. and Urtasun, R. (n.d.). *RadarNet: Exploiting Radar for Robust Perception of Dynamic Objects*. [online] Available at: <https://arxiv.org/pdf/2007.14366.pdf>.

Zhou, T., Yang, M., Jiang, K., Wong, H.T. and Yang, D. (2020). MMW Radar-Based Technologies in Autonomous Driving: A Review. *Sensors*, [online] 20(24), pp.7283–7283. doi:<https://doi.org/10.3390/s20247283>.

Zhou, Y. and Yue, Y. (2022). *Radar Signal Processing Fundamentals*. [online] Encyclopedia.pub. Available at: <https://encyclopedia.pub/entry/23781> [Accessed 23 Oct. 2023].

Joel Markus Vaz and Balaji, S. (2021). Convolutional neural networks (CNNs): concepts and applications in pharmacogenomics. *Molecular Diversity*, [online] 25(3), pp.1569–1584. doi:<https://doi.org/10.1007/s11030-021-10225-3>.

Deng, J., Dong, W., Socher, R., Li, L., Li, K. and Li, F. (2009). ImageNet: A large-scale hierarchical image database. *2009 IEEE Conference on Computer Vision and Pattern Recognition*. [online] doi:<https://doi.org/10.1109/cvpr.2009.5206848>.

Hironobu Fujiyoshi, Hirakawa, T. and Yamashita, T. (2019). Deep learning-based image recognition for autonomous driving. *IATSS Research*, [online] 43(4), pp.244–252. doi:<https://doi.org/10.1016/j.iatssr.2019.11.008>.

Balasubramaniam, A. and Pasricha, S. (n.d.). *Object Detection in Autonomous Vehicles: Status and Open Challenges*. [online] Available at: <https://arxiv.org/ftp/arxiv/papers/2201/2201.07706.pdf#:~:text=Object%20detection%20consists%20of%20two>.

Nitin Kushwaha (2023). *A Brief History of the Evolution of Image Classification*. [online] Medium. Available at: <https://python.plainenglish.io/a-brief-history-of-the-evolution-of-image-classification-402c63baf50>

Deepchecks. (2021). *What is VGGNet | Deepchecks*. [online] Available at: <https://deepchecks.com/glossary/vggnet/#:~:text=Its%20object%20recognition%20method%20developed,dataset%20by%20a%20wide%20margin>.

Great Learning Team (2020). *Introduction to Resnet or Residual Network*. [online] Great Learning Blog: Free Resources what Matters to shape your Career! Available at: <https://www.mygreatlearning.com/blog/resnet/#:~:text=ResNet%2C%20short%20for%20Residual%20Network,Residual%20Learning%20for%20Image%20Recognition%E2%80%9D>.

Paperswithcode.com. (2020). *Papers with Code - DenseNet Explained*. [online] Available at: [https://paperswithcode.com/method/densenet/#:~:text=A%20DenseNet%20is%20a%20type,sizes\)%20directly%20with%20each%20other](https://paperswithcode.com/method/densenet/#:~:text=A%20DenseNet%20is%20a%20type,sizes)%20directly%20with%20each%20other).

Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J. and Houlsby, N. (2020). *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*. [online] arXiv.org. Available at: <https://arxiv.org/abs/2010.11929>.

GitHub. (2024). *tudelft-iv/view-of-delft-dataset*: This repository shares the documentation and development kit of the View of Delft automotive dataset. [online] Available at: <https://github.com/tudelft-iv/view-of-delft-dataset?tab=readme-ov-file#introduction>

kili-website. (2023). *Mean Average Precision (mAP): A Complete Guide*. [online] Available at: <https://kili-technology.com/data-labeling/machine-learning/mean-average-precision-map-a-complete-guide>

GitHub. (2024). *OpenPCDet/docs/multiple_models_demo.png at master · open-mmlab/OpenPCDet*. [online] Available at: https://github.com/open-mmlab/OpenPCDet/blob/master/docs/multiple_models_demo.png

Guan, S., Wan, C. and Jiang, Y. (2023). *AV PV-RCNN: Improving 3D Object Detection with Adaptive Deformation and VectorPool Aggregation*. [online] ResearchGate. Available at: https://www.researchgate.net/publication/371321106_AV_PV-RCNN_Improving_3D_Object_Detection_with_Adaptive_Deformation_and_VectorPool_Aggregation

Aha (2024). Agile Software Development - Agile Methodology Explained. [online] Wwww.aha.io. Available at: <https://www.aha.io/roadmapping/guide/agile/agile-software-development>